

Chapter 2 — Methodology and Empirical

Wind-Speed Modelling

Continuous-Time Modelling of Wind Speed for Derivatives Pricing and Risk Management

This chapter develops the wind-speed modelling pipeline whose motivation and literature grounding were set out in Chapter 1: a five-stage sequence running from deseasonalisation through discrete-time autoregression and conditional volatility to a continuous-time embedding and, finally, a theoretical stochastic-volatility extension. Each stage is applied identically to both datasets described in Section 1.4: Dataset A (Bologna, 10 m, 2015–2018), on which the pipeline is developed and validated, and Dataset B (Nordfriesland, 100 m, 2015–2024), which carries the thesis’s headline results. Each section below presents the two datasets side by side, so that methodological choices can be compared directly against the data they are applied to. Numerical results quoted throughout are frozen outputs of the underlying phase notebooks; none are re-estimated here.

1 Deseasonalisation and the Box–Cox Transformation

Raw daily wind speed is a positive, right-skewed, seasonally heteroskedastic quantity, and every subsequent stage of this pipeline (the AR(4) mean model of Section 2, the GARCH(1,1) volatility layer of Section 3, and the CAR(4) continuous-time embedding of Section 4) presupposes a series that is approximately Gaussian, zero-mean, and unit-variance. The first stage of the pipeline is therefore a sequence of four operations applied to raw wind speed $W(t)$: characterising its marginal distribution, applying a Box–Cox power transformation, removing a deterministic seasonal mean $\Lambda(t)$ via a truncated Fourier series, and removing a deterministic seasonal variance $\sigma^2(t)$ via a second Fourier fit. The output is the standardised series $\tilde{W}(t) = [W^{(\lambda)}(t) - \Lambda(t)]/\sigma(t)$, which enters the AR(4) model of Section 2.

1.1 Marginal Distribution

The two marginal distributions fitted throughout this section are the Weibull and the lognormal,

$$f_W(v; k, c) = \frac{k}{c} \left(\frac{v}{c}\right)^{k-1} \exp \left[-\left(\frac{v}{c}\right)^k \right], \quad v \geq 0, \quad (1)$$

$$f_L(v; m, s) = \frac{1}{vs\sqrt{2\pi}} \exp \left[-\frac{(\ln v - m)^2}{2s^2} \right], \quad v > 0, \quad (2)$$

both fitted by maximum likelihood and compared by Akaike information criterion. Table 1 reports the resulting fits for both datasets. Tuller and Brett [1984] provide the physical justification for the Weibull distribution reviewed in Section 1.3.1: under approximate bivariate-normal isotropy of the horizontal wind-velocity components, wind speed is Weibull-distributed with shape parameter near two. Dataset B, measured at 100 m in the free atmosphere above the surface boundary layer, conforms to this picture: the Weibull distribution ($\hat{k} = 2.67$) beats the lognormal alternative on both AIC and BIC, consistent with Tuller and Brett [1984] and with the case-study evidence of Garcia et al. [1998]. Dataset A does not: at 10 m, where surface friction and diurnal boundary-layer cycling distort the isotropy assumption underlying the Weibull derivation, the lognormal distribution wins decisively (log-likelihood $-1,730.6$ versus $-1,960.5$ for the Weibull MLE fit, $\Delta\text{AIC} \approx 460$). This is the same qualitative finding already noted in Section 1.4.1; it is repeated here because it is the reason the two datasets receive different treatment in the Box–Cox step that follows.

Table 1: Marginal distribution fits, both datasets (maximum likelihood).

	Dataset A	Dataset B
Weibull ($\hat{k}, \hat{c}$)	2.6368, 2.7729 m/s	2.67, 9.09 m/s
Lognormal ($\hat{m}, \hat{s}$)	0.8423, 0.3405	2.00, 0.43
Log-likelihood, Weibull	$-1,960.51$	n/a*
Log-likelihood, Lognormal	$-1,730.56$	n/a*
AIC, Weibull / Lognormal	3,925.01 / 3,465.12	18,758.3 / 18,803.1
Better fit	Lognormal ($\Delta\text{AIC} \approx 460$)	Weibull ($\Delta\text{AIC} \approx 45$)

*Dataset A's log-likelihoods are reported in the source analysis; Dataset B's own report gives AIC/BIC directly without a separately tabulated log-likelihood column.

1.2 Box–Cox Transformation

The Box–Cox family $W^{(\lambda)}(t) = [W(t)^\lambda - 1]/\lambda$ (with $\ln W(t)$ for $\lambda = 0$) is fitted by maximising the profile log-likelihood under a normality assumption. For Dataset A, the likelihood-maximising value is $\lambda_{\text{opt}} = -0.450$, achieving near-perfect Gaussianisation (skewness -0.005 , excess kurtosis $+0.128$, Jarque–Bera $p = 0.606$). For Dataset B, the optimum is $\lambda_{\text{opt}} = 0.404$, a much less extreme transformation, consistent with the higher hub-height mean wind speed producing a less severely skewed raw distribution. In both cases, the pricing formulas developed in Section 4 and used throughout Chapter 3 require $1/\lambda$ to be a positive integer, so the likelihood optimum cannot be used directly. The




Figure 1: Raw wind speed distribution and Box–Cox transformation, both datasets. Each panel is arranged in four rows, read top to bottom. **Row 1:** the raw daily wind speed time series, with the sample mean marked by a dashed line. **Row 2:** three diagnostics for the raw series – a histogram with the fitted Weibull and lognormal densities overlaid, a Weibull probability-paper linearisation, and a quantile–quantile plot against the normal distribution, which shows the characteristic S-curve deviation of a right-skewed, heavy-tailed series. **Row 3:** histogram and quantile–quantile plot of the series after the likelihood-optimal Box–Cox transformation λ_{opt} . **Row 4:** histogram and quantile–quantile plot after the tractable Box–Cox value actually adopted ($\lambda = 0.1$ for Dataset A, $\lambda = 0.5$ for Dataset B), showing the residual departure from normality left by the tractability constraint discussed in Section 1.2.

requirement that $1/\lambda \in \mathbb{Z}^+$ is dictated by the analytical pricing formulas of Chapter 3. Pricing a forward contract requires computing the expected future wind speed in its original physical measure. Inverting the Box–Cox transformation requires taking the $(1/\lambda)$ -th moment of a normal random variable, denoted $M_{1/\lambda}$ [Benth and Šaltytė Benth, 2009]. If $1/\lambda$ were fractional, this expectation would lack a closed-form solution and risk generating complex numbers from negative intermediate states, destroying the model’s analytical tractability. Following Benth and Šaltytė Benth [2009], who adopt $\lambda \in \{0.1, 0.2\}$ for the same CAR(p) continuous-time embedding, $\lambda = 0.1$ is adopted for Dataset A and $\lambda = 0.5$ for Dataset B (Table 2). The rounding cost is asymmetric: Dataset A’s tractable value sits 0.35 units from its optimum and leaves a residual right skew (skewness 0.601, Jarque–Bera $p \approx 0$), while Dataset B’s tractable value sits only 0.10 units from its optimum and retains near-Gaussian shape (skewness 0.081, KS $p = 0.091$).

Table 2: Box–Cox transformation: optimal versus tractable λ .

	Dataset A	Dataset B
λ_{opt} (likelihood-maximising)	−0.450	0.404
Gaussianisation at λ_{opt} (skew, JB p)	−0.005, 0.606	−0.020, 0.0002
λ adopted (tractable, $1/\lambda \in \mathbb{Z}^+$)	0.1 ($1/\lambda = 10$)	0.5 ($1/\lambda = 2$)
Residual shape at adopted λ (skew, JB p)	0.601, ≈ 0	0.081, < 0.001

1.3 Seasonal Mean and Seasonal Variance

Both the seasonal mean $\Lambda(t)$ and the seasonal variance $\sigma^2(t)$ are fitted as two-harmonic truncated Fourier series in calendar day-of-year, $\Lambda(t) = a_0 + \sum_{k=1}^2 [a_k \cos(2\pi kt/365) + b_k \sin(2\pi kt/365)]$ by ordinary least squares, and $\sigma^2(t) = c_0 + \sum_{k=1}^2 c_k \cos(2\pi kt/365)$ (cosine-only, reflecting the approximate symmetry of the variance cycle about its annual peak) fitted to the per-calendar-day empirical variance of the residuals $\varepsilon(t) = W^{(l)}(t) - \Lambda(t)$. Table 3 reports both fits. The annual harmonic dominates the seasonal mean in both datasets (b_1 significant, a_1 not, for Dataset A; a_1 significant, b_2 not, for Dataset B), and the fitted R^2 is modest in both cases (0.126 and 0.103): most of the variance in daily wind speed is weather-driven, not calendar-driven, a finding consistent across both the four-year proxy sample and the ten-year reference sample. The seasonal-variance fit is more informative for the purposes of this thesis, since its intercept C_0 is the Fourier constant that enters the variance-swap pricing formula $K_{\text{var}}^B = h \times C_0$ used throughout Chapter 3: $C_0 = 0.13138$ for Dataset A versus $C_0 \approx 1.163$ for Dataset B, an 8.8-fold gap that Chapter 3's Value-of-Information analysis identifies as the dominant driver of the pricing error attributable to using the geographic proxy. Both datasets place peak seasonal variance in winter, though the relative modulation differs sharply: Dataset A's fitted $\sigma^2(t)$ ranges over $[0.032, 0.224]$, a winter-to-summer ratio of approximately 6.9 to 1, versus Dataset B's much milder range of $[0.944, 1.463]$, a peak-to-trough increase of about 55%.

Table 3: Seasonal mean and seasonal variance Fourier fits, both datasets.

	Dataset A	Dataset B
<i>Seasonal mean $\Lambda(t)$</i>		
a_0	0.88513 (***)	3.564 (***)
a_1 (annual cos)	−0.00544 (n.s.)	0.510 (***)
b_1 (annual sin)	0.16381 (***)	0.083 (**)
R^2	0.126	0.103
<i>Seasonal variance $\sigma^2(t)$</i>		
C_0 (mean level)	0.13138 (***)	1.16277 (***)
c_1 (annual cos)	0.09579 (***)	0.25933 (***)
c_2 (semi-annual cos)	−0.00332 (n.s.)	0.04077 (n.s.)
R^2	0.226	0.116
$\sigma^2(t)$ range	$[0.032, 0.224]$	$[0.944, 1.463]$

1.4 The Standardised Series $\tilde{W}(t)$

The standardised series $\tilde{W}(t) = \varepsilon(t)/\sigma(t)$ is, by construction, close to zero-mean and unit-variance in both datasets (Dataset A: mean 0.001, std. 0.967; Dataset B: mean −0.00003, std. 1.005), confirming that the centring and scaling steps were executed correctly. Neither series is yet white noise, however. The Ljung–Box statistic tests the

joint null hypothesis that the first m autocorrelations of a series are zero,

$$LBQ(m) = n(n+2) \sum_{k=1}^m \frac{\hat{\rho}(k)^2}{n-k} \sim \chi^2(m) \quad \text{under } H_0, \quad (3)$$

where $\hat{\rho}(k)$ is the sample autocorrelation at lag k ; applied to a series' squared values, the same statistic tests for conditional heteroskedasticity rather than serial correlation in levels. The test rejects the null of no serial correlation at every tested lag (5, 10, 20) for $\tilde{W}(t)$ itself in both datasets ($p < 0.001$), and the same test applied to $\tilde{W}(t)^2$ also rejects at every lag in both datasets ($p < 0.001$), indicating conditional heteroskedasticity beyond the deterministic seasonal variance already removed. These two findings are the direct motivation for the AR(4) mean model of Section 2 and the GARCH(1,1) volatility model of Section 3 respectively, and the pattern is identical across both datasets despite their very different sample lengths and measurement heights. This is evidence that the residual dynamics being modelled are a genuine feature of daily wind speed, not an artefact of one particular site or sample.




Figure 2: Seasonal decomposition and standardisation, both datasets. Each panel is arranged in three rows. **Row 1:** the transformed wind speed $W^{(l)}(t)$ with the fitted seasonal mean $\Lambda(t)$ overlaid. **Row 2:** three panels – the deseasonalised residuals $\varepsilon(t) = W^{(l)}(t) - \Lambda(t)$; the empirical per-calendar-day variance (points) with the fitted Fourier seasonal-variance curve $\sigma^2(t)$ (line), showing the winter peak / summer trough pattern common to both datasets; and, third, an autocorrelation plot of the seasonal standard deviation curve across the calendar year, a supplementary panel not otherwise discussed in the main text. **Row 3:** the standardised series $\tilde{W}(t)$, its histogram against a normal density, and a quantile–quantile plot, illustrating the residual serial dependence and non-normality quantified in Section 1.4.

2 Conditional Mean: AR(4)

The Ljung–Box evidence of Section 1.4 establishes that $\tilde{W}(t)$ retains serial dependence. This section fits an autoregressive model to capture it.

2.1 Order Selection

Both datasets exhibit a similar autocorrelation profile: a partial autocorrelation function (PACF) with clearly significant spikes at lags 1–3 and a markedly weaker, borderline value at lag 4 (Dataset A: $\text{PACF}_4 = 0.0311$, just inside the 95% band; Dataset B: $\text{PACF}_4 = -0.0070$, effectively zero), with no significant structure beyond lag 4 in either series. An AIC/BIC grid search over $\text{AR}(p)$, $p = 1, \dots, 6$, in fact selects $\text{AR}(3)$ as the minimum-AIC model for *both* datasets (Dataset A: $\text{AIC} = 3,432.24$ at $p = 3$ versus $3,432.83$ at $p = 4$, a difference of 0.59 units; Dataset B: $\text{AIC} = 9,342.24$ at $p = 3$ versus $9,344.04$ at $p = 4$, a difference of 1.8 units), and in both cases the gap falls within the conventional two-point AIC band below which two models are considered statistically indistinguishable.

$\text{AR}(4)$ is nonetheless retained as the operating specification in both datasets, on structural rather than statistical grounds: the continuous-time $\text{CAR}(4)$ embedding of Section 4 requires a companion matrix of dimension four, and the Euler-discretisation mapping from discrete AR coefficients $\{\beta_1, \dots, \beta_4\}$ to continuous CAR coefficients $\{\alpha_1, \dots, \alpha_4\}$, following Benth and Šaltytė Benth [2009], does not generalise to a three-lag or mixed ARMA specification. This is a deliberate trade-off, stated candidly here rather than left implicit: a fractionally-smaller-AIC three-lag model is set aside in favour of the four-lag model that the pricing framework of Chapter 3 actually requires, on the grounds that the statistical cost of doing so is, in both datasets, within the range normally treated as noise.

2.2 AR(4) Coefficients

Table 4 reports the fitted $\text{AR}(4)$ coefficients for both datasets, estimated on the zero-mean series $\tilde{W}(t)$ without an intercept. Both datasets show the coefficient magnitude declining with lag order and a positive, dominant β_1 , but the two series differ in the sign of β_4 (Dataset A: $+0.0310$; Dataset B: -0.0074) and in the statistical significance of the higher lags: Dataset A's β_2 and β_3 are both significant at the 1% level and β_4 is not significant ($p = 0.208$), while Dataset B's β_3 is only marginally significant ($p = 0.090$) and β_4 is clearly insignificant ($p = 0.660$). Neither discrepancy affects the modelling decision, since $\text{AR}(4)$ is retained structurally regardless of the marginal significance of its highest lag, but both are reported here rather than smoothed over, since a coefficient that changes sign between the two datasets is a legitimate empirical finding, not a data-quality concern.

Both coefficient sets can be read against the New York daily wind-speed $\text{AR}(4)$ fit reported by Alexandridis and Zapranis [2013]: $\beta_1 = 0.362$, $\beta_2 = -0.100$, $\beta_3 = 0.027$, $\beta_4 = 0.022$. Dataset A's coefficients are larger in magnitude than the New York benchmark at every lag. Dataset B's are larger at β_1 (0.532 versus 0.362) and β_3 (0.032 versus 0.027), but smaller at β_2 (0.076 versus 0.100 in magnitude) and β_4 (0.007 versus 0.022), so the comparison holds cleanly for the proxy dataset but only partially for the reference dataset. The dominant lag-1 coefficient, which carries most of the

Table 4: AR(4) coefficients, both datasets (OLS/MLE on $\tilde{W}(t)$, no intercept).

Coefficient	Dataset A		Dataset B	
	Estimate	<i>p</i> -value	Estimate	<i>p</i> -value
β_1	0.6299	< 0.001	0.5324	< 0.001
β_2	−0.1443	< 0.001	−0.0763	< 0.001
β_3	0.0926	0.001	0.0324	0.090
β_4	0.0310	0.208	−0.0074	0.660

autoregressive structure in both series, is larger at both European sites than at the New York site, and this is attributed, in the underlying phase analysis, to a stronger lag-1 sample autocorrelation at both European sites (0.579 for Dataset A) than at the New York site (≈ 0.36), consistent with the sites’ differing synoptic exposure rather than with any difference in modelling approach: the AR(4) specification itself is identical to the one used in the benchmark study.

For Dataset A, the characteristic roots of the fitted AR(4) polynomial have moduli $\{0.702, 0.470, 0.470, 0.200\}$, all inside the unit circle, confirming stationarity; the dominant root implies a mean-reversion half-life for the level of wind-speed anomalies of approximately 2.0 days. This is a materially faster reversion than the volatility half-life derived from the GARCH model in Section 3 (41.6 and 38.3 days for Datasets A and B respectively): the *level* of a wind-speed anomaly washes out within days, while the *volatility regime* it occurred in persists for weeks, a distinction that recurs when the two dynamics are combined in Section 4.

2.3 Residual Diagnostics

Table 5 summarises the Ljung–Box tests on the AR(4) residuals $\hat{\varepsilon}(t)$ and on their squares. In both datasets, the serial-correlation test on levels passes comfortably at every lag ($p > 0.5$ throughout), confirming that AR(4) has removed the linear dependence identified in Section 1.4. In both datasets, the test on squared residuals rejects the null of no remaining structure at conventional significance. This is corroborated by the Engle [1982] Lagrange-multiplier test for autoregressive conditional heteroskedasticity, which regresses squared residuals on their own lags, $\hat{\varepsilon}(t)^2 = \gamma_0 + \sum_{k=1}^q \gamma_k \hat{\varepsilon}(t-k)^2 + u(t)$, and tests the joint significance of $\gamma_1, \dots, \gamma_q$ via $LM(q) = nR^2 \sim \chi^2(q)$ under the null of no ARCH effects; the entry-residual ARCH-LM test is significant in both datasets (Dataset A: $LM(5) = 18.80$, $p = 0.0021$; Dataset B: $LM(5) = 16.58$, $p = 0.0054$). This is the expected and desired outcome at this stage: the AR(4) mean equation is adequate, and the conditional heteroskedasticity that remains is precisely what the GARCH(1,1) layer of Section 3 is designed to capture.

2.4 Forecast Accuracy versus Persistence

Table 6 reports one-step-ahead out-of-sample forecast accuracy against a persistence benchmark $\hat{\tilde{W}}(t) = \tilde{W}(t-1)$, using a walk-forward evaluation held out from the estima-

Table 5: AR(4) residual diagnostics, both datasets.

Test	Dataset A	Dataset B
Ljung–Box $\hat{\varepsilon}(t)$, lag 20 (p)	0.506	0.511
Ljung–Box $\hat{\varepsilon}(t)^2$, lag 20 (p)	0.0033	0.2724*
ARCH-LM(5) on entry residuals (p)	0.0021	0.0054

analysis is internally inconsistent on its squared-residual diagnostics: the narrative text reports that the Ljung–Box test rejects homoskedasticity at lags 5 and 10 ($p = 0.0048$ and 0.0419), while the lag-20 value tabulated there is 0.2724 (a pass). The cell above reports the lag-20 value from the source table. Note that lags 5 and 10 rejections are the substantive finding motivating GARCH, regardless of the lag-20 outcome, and the ARCH-LM entry-residuals test (final row) is independently significant in both datasets.

tion sample. AR(4) reduces RMSE relative to persistence by 12.2% in Dataset A and 13.8% in Dataset B, and a Diebold–Mariano test [Diebold and Mariano, 1995] rejects the null of equal predictive accuracy in both cases (Dataset A: $DM = -3.681$, $p = 0.0002$; Dataset B: $DM = -7.019$, $p < 0.001$). These are the frozen benchmark figures for the two datasets and are reused, without re-estimation, in the cross-dataset validation exercise of Chapter 3; a separate, later validation split reported there for Dataset B under the same Diebold–Mariano procedure (Section 3.2) uses a different train/test partition and yields a similar but not numerically identical improvement (13.6%), a difference attributable entirely to the split rather than to any change in the fitted model, and reported explicitly as such where it recurs.

Table 6: AR(4) vs. persistence, one-step-ahead out-of-sample RMSE.

	Dataset A		Dataset B	
	RMSE	MAE	RMSE	MAE
Persistence	0.6856	0.5214	0.9955	0.8054
AR(4)	0.6017	0.4706	0.8580	0.6950
Improvement	12.2%	—	13.8%	—

2.5 Long Memory

A natural question, given the literature reviewed in Section 1.3.2, is whether the residual dependence in $\tilde{W}(t)$ decays exponentially, as a finite-order AR model assumes, or hyperbolically, as fractional integration would imply. The Geweke–Porter–Hudak (GPH) log-periodogram estimator regresses the log-periodogram of the series on the log of the frequency, $\ln I(\omega_j) = c - d \ln [4 \sin^2(\omega_j/2)] + u_j$, over the $m = n^{0.6}$ lowest Fourier frequencies ω_j , and tests $H_0 : d = 0$ against fractional integration $d \neq 0$. It gives a fractional-differencing parameter of $\hat{d} = 0.119$ ($p = 0.208$) for Dataset A and $\hat{d} = -0.115$ ($p = 0.094$) for Dataset B. Applying the same $p \leq 0.10$ decision rule to both datasets, Dataset A cleanly fails to reject $d = 0$ and no ARFIMA model is fitted; Dataset B falls just inside the threshold, and an ARFIMA(4, d , 0) specification was estimated




Figure 3: AR(4) identification and residual diagnostics, both datasets. Each panel is arranged in three rows. **Row 1:** three panels – the autocorrelation function of $\tilde{W}(t)$; its partial autocorrelation function, showing the sharp PACF cutoff after lag 4 discussed in Section 2.1; and the AIC grid search across candidate AR orders. **Row 2:** the time series of the AR(4) residuals $\hat{\varepsilon}(t)$. **Row 3:** three further panels, confirming that AR(4) has removed the linear serial dependence identified in Section 1.4 while conditional heteroskedasticity remains – an autocorrelation-type diagnostic of $\hat{\varepsilon}(t)$; the corresponding diagnostic applied to $\hat{\varepsilon}(t)^2$, which stays visibly elevated and motivates the GARCH model of Section 3; and a histogram of $\hat{\varepsilon}(t)$ against a normal density.

as a robustness check, reporting a substantially lower AIC than the AR(4) benchmark.¹ Read cautiously, given that footnote, the honest summary is that Dataset A shows no evidence of long memory at daily frequency, while Dataset B shows marginal, borderline evidence that does not decisively overturn the finite-order AR(4) specification adopted here; Haslett and Raftery [1989] and Kavasseri and Seetharaman [2009] both find long memory to be primarily a feature of hourly or sub-daily wind data, which is consistent with a weak or borderline signal surviving daily aggregation in at most one of the two datasets used here. AR(4) is retained as the operating mean model for both datasets, for the same structural reason given in Section 2.1.

3 Conditional Volatility: GARCH(1,1)

Section 2.3 established that AR(4) residuals remain conditionally heteroskedastic in both datasets. This section models that heteroskedasticity directly, completing the two-component variance decomposition $\sigma_{\text{total}}^2(t) = \sigma_{\text{seasonal}}^2(t) \times h(t)$ that combines the deterministic Fourier seasonal variance of Section 1.3 with a GARCH(1, 1) conditional-

¹The exact value of d used in Dataset B’s own ARFIMA robustness fit could not be reconciled with the GPH hypothesis-test estimate reported immediately above it in the same source document ($\hat{d} = -0.115$ for the test versus a value of 0.010 referenced for the ARFIMA fit itself); the source document’s own Hurst-exponent figure caption also states that its R/S estimate “confirms the absence of strong long-range dependence,” in tension with the surrounding text’s conclusion that long memory “is present.” These are inconsistencies internal to the Dataset B source analysis, not something resolved by re-estimation here. Due to this internal inconsistency regarding the exact value of d used in the ARFIMA fit, the GPH hypothesis-test estimate ($\hat{d} = -0.115$) is prioritised in the main text above, rather than the value referenced for the ARFIMA fit itself.

variance multiplier $h(t)$, following Tol [1997] and Bollerslev [1986].

3.1 GARCH(1,1) Estimates

Table 7 reports the fitted GARCH(1,1) parameters for both datasets. In both cases β dominates α by more than an order of magnitude (Dataset A: $\beta/\alpha \approx 31$; Dataset B: $\beta/\alpha \approx 165$), meaning that conditional variance is overwhelmingly carried forward from its own past level rather than driven by fresh shocks. Tol [1997] characterises this pattern as typical of daily meteorological series, where volatility is governed by slow-moving synoptic weather regimes rather than the sharp, episodic jumps observed in financial return data. Dataset B's β is both larger and more precisely estimated than Dataset A's, while its α is smaller and, unlike Dataset A's, not statistically significant on its own ($p = 0.210$); the constant ω is likewise insignificant for Dataset B ($p = 0.348$) but marginally significant for Dataset A ($p = 0.060$). Neither insignificant coefficient is dropped from either model, since the two-component variance decomposition requires all three GARCH parameters to reconstruct $h(t)$, and persistence, half-life, and the implied unconditional variance are computed from the fitted values regardless of the individual coefficients' marginal significance.

Table 7: GARCH(1,1) estimates, both datasets.

Parameter	Dataset A		Dataset B	
	Estimate	p -value	Estimate	p -value
ω	0.01004	0.060	0.01348	0.348
α	0.03105	< 0.001	0.00590	0.210
β	0.95242	< 0.001	0.97618	< 0.001
Persistence $\alpha + \beta$	0.9835		0.9821	
Half-life (days)	41.6		38.3	
Implied unconditional variance	0.608		0.752	
$h(t_0)$ (end of sample)	0.4137		0.7511	

Both models are covariance-stationary ($\alpha + \beta < 1$ in both cases). The volatility half-life (the time for a unit shock to conditional variance to decay to half its initial size, $\ln 2 / \ln[1/(\alpha + \beta)]$) is 41.6 days for Dataset A and 38.3 days for Dataset B, an order of magnitude longer than the roughly two-day mean-reversion half-life of the wind-speed *level* found in Section 2.2. Table 8 traces the resulting impulse-response decay $(\alpha + \beta)^{\text{lag}}$ explicitly: a unit shock to conditional variance still retains roughly 70% of its initial effect after 20 days in both datasets (71.6% for Dataset A, 69.7% for Dataset B), and does not fall below one-quarter of its initial size until roughly three months out.

The distinction is physically meaningful and consistent across both datasets: an individual wind-speed anomaly washes out within days, but the volatility *regime* it belongs to persists for weeks, consistent with the multi-week persistence of large-scale atmospheric circulation patterns such as blocking events and North Atlantic Oscillation phase.

Table 8: GARCH(1,1) impulse-response decay, $(\alpha + \beta)^{\text{lag}}$, both datasets.

Lag (days)	Dataset A	Dataset B
0	100.0%	100.0%
5	92.0%	91.4%
10	84.6%	83.5%
20	71.6%	69.7%
30	—	58.4%
38 (Dataset B half-life)	—	50.0%
42 (Dataset A half-life)	50.0%	—
90	22.3%	—

3.2 Two-Component Variance Decomposition

The key numeric output of this stage for the pricing work of Chapter 3 is the pair $(h(t_0), C_0)$: the GARCH multiplier evaluated at the end of the sample, and the Fourier seasonal-variance constant from Section 1.3. Chapter 3’s synthetic variance swap prices under a piecewise-constant approximation, following Tol [1997], in which $h(t)$ is frozen at $h(t_0)$ for the life of the contract: $h(t_0) = 0.4137$ for Dataset A and $h(t_0) = 0.7511$ for Dataset B. Combined with $C_0 = 0.13138$ and $C_0 \approx 1.163$ respectively (Section 1.3), this gives $\sigma_{\text{total}}^2(t_0) = 0.0926$ for Dataset A, an exact reproduction of the frozen anchor quoted in the project specification, confirming internal consistency between the Phase 1 seasonal-variance fit and the Phase 3 GARCH fit.

The seasonal pattern of $h(t)$ itself differs qualitatively between the two datasets. In Dataset A, $h(t)$ is markedly higher in late spring and early summer (monthly mean up to 0.768 in June) than in autumn (as low as 0.465 in October), a pattern that runs in the opposite direction to the seasonal variance $\sigma_{\text{seasonal}}^2(t)$, which peaks in winter; the two effects partially offset each other in the total variance $\sigma_{\text{total}}^2(t)$, so that pricing from the seasonal component alone, ignoring $h(t)$, would overstate summer volatility and understate the GARCH contribution in the shoulder months. In Dataset B, by contrast, $h(t)$ shows only a mild seasonal range (monthly means between 0.734 and 0.770), and most of its time variation is carried by shock persistence rather than calendar effects; total variance in Dataset B therefore tracks the seasonal component more directly than in Dataset A.

3.3 Standardised Residual Diagnostics

Table 9 reports diagnostics on the standardised residuals $z(t) = \hat{\varepsilon}(t)/\sqrt{h(t)}$. Both datasets pass the Ljung–Box test for remaining serial correlation and remaining ARCH effects on $z(t)$ and $z(t)^2$ at conventional significance, confirming that GARCH(1,1) is an adequate specification for the conditional-variance dynamics in both cases. The two datasets diverge, however, in the shape of their standardised residuals: Dataset A’s $z(t)$ is leptokurtic (excess kurtosis +1.12), with a heavy right tail driven substantially by a single extreme observation (28 June 2017, $z = +5.79$), while Dataset B’s is mildly platykurtic (excess kurtosis −0.28), thinner-tailed than Gaussian. Neither departure

is treated as invalidating the GARCH(1,1) specification, since both series pass the joint white-noise tests that matter for the volatility model itself, but the residual non-normality, in whichever direction, motivates the heavy-tailed Normal-Inverse-Gaussian treatment introduced in Section 4.4. One further asymmetry is worth flagging directly rather than smoothing over: an ARCH-LM test on Dataset B’s standardised residuals is marginally significant at 5% ($p = 0.0185$), where the equivalent test on Dataset A’s standardised residuals passes comfortably ($p = 0.821$); Dataset B’s own summary narrative characterises its residuals as passing validation cleanly, a characterisation that this specific test result only partially supports. This is not an artefact of a different evaluation window: the same lag-5 statistic and p -value recur, unchanged, in the later cross-dataset validation of Chapter 3, which reaches the same qualified verdict on the same evidence.

Table 9: GARCH(1,1) standardised-residual diagnostics, both datasets.

Test	Dataset A	Dataset B
Excess kurtosis of $z(t)$	+1.12 (leptokurtic)	−0.28 (platykurtic)
Ljung–Box $z(t)$, lag 20 (p)	0.765	0.503
Ljung–Box $z(t)^2$, lag 20 (p)	0.988	0.283
ARCH-LM(5) on $z(t)$ (p)	0.821	0.0185




Figure 4: GARCH(1,1) conditional-volatility decomposition, both datasets. Each panel is arranged in four rows. **Row 1:** the AR(4) residuals with GARCH-adaptive 95% confidence bands ($\pm 1.96\sqrt{h(t)}$) plotted against fixed-width bands, visibly widening around high-volatility episodes. **Row 2:** the fitted conditional-variance path $h(t)$ against its unconditional long-run level. **Row 3:** the two-component total volatility $\sigma_{\text{total}}(t) = \sqrt{\sigma_{\text{seasonal}}^2(t) h(t)}$ against the Phase 1 seasonal baseline. **Row 4:** three panels – a histogram of the standardised residuals $z(t)$ against the normal distribution; a quantile–quantile plot of $z(t)$ against the normal distribution; and a monthly bar-chart comparison of the seasonal-variance and GARCH components, illustrating the dataset-dependent degree to which the two effects offset each other, discussed in Section 3.2.

3.4 Cross-Dataset Comparison

Dataset B’s own analysis draws an explicit comparison back to Dataset A, attributing its lower shock-sensitivity ($\alpha = 0.0059$ versus 0.0311) and higher persistence-to-shock ratio ($\beta/\alpha = 165$ versus 31) to more stable boundary-layer mixing at 100 m hub height relative to the 10 m surface layer, over a decade-long record rather than a four-year one. One phrase in that comparison is worth a precision note: Dataset B’s total persistence $\alpha + \beta = 0.9821$ is in fact marginally *lower* than Dataset A’s 0.9835 , so the claim that Dataset B is “more persistent” rests on the β -to- α decomposition, not on $\alpha + \beta$ itself; this thesis reports the comparison on that more precise basis in Table 7 rather than repeating the aggregate persistence framing.

4 Continuous-Time Embedding: CAR(4)

The AR(4) model of Section 2 is fit on a daily grid, but a derivative contract must in general be priced at an arbitrary maturity that does not fall on that grid. This section embeds the discrete AR(4) model into a continuous-time CAR(4) state-space representation, following Benth and Šaltytė Benth [2009] and the general Lévy-driven CARMA theory of Brockwell and Marquardt [2005].

4.1 Euler Discretisation and the Companion Matrix

The mapping between the two representations is not an independent fit: it is the algebraic consequence of a single modelling choice, namely that the daily AR(4) recursion is itself what results from discretising the continuous-time CAR(4) stochastic differential equation, $dX(t) = A X(t) dt + e_4 \sigma(t) dB(t)$, by a first-order (Euler) finite-difference approximation taken at a one-day step size. Writing that finite-difference approximation out and matching it term by term against the discrete AR(4) equation yields four algebraic identities expressing the continuous-time coefficients $\{\alpha_1, \dots, \alpha_4\}$ directly in terms of the discrete $\{\beta_1, \dots, \beta_4\}$, given by Benth and Šaltytė Benth [2009]: $\alpha_1 = 4 - \beta_1$, $\alpha_2 = 3\alpha_1 - \beta_2 - 6$, $\alpha_3 = 4 + 2\alpha_2 - \beta_3 - 3\alpha_1$, $\alpha_4 = \alpha_3 - \beta_4 - \alpha_2 + \alpha_1 - 1$. Because the map is derived algebraically rather than fitted, it is invertible by construction, and for Dataset A this is verified directly: substituting the fitted α coefficients back through the inverse relations reproduces the original β coefficients to machine precision ($\max |\beta_{\text{orig}} - \beta_{\text{reconstructed}}| = 7.77 \times 10^{-16}$), confirming that no information is lost in either direction of the mapping. The resulting α coefficients populate the last row of the companion matrix

$$A = \begin{pmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ -\alpha_4 & -\alpha_3 & -\alpha_2 & -\alpha_1 \end{pmatrix}, \quad (4)$$

following the general companion-matrix construction of Brockwell and Marquardt [2005]. Table 10 reports the fitted coefficients for both datasets: Dataset B’s coeffi-

cients are larger than Dataset A's at every lag, and the gap widens with lag order, from a 2.9% difference at α_1 to 5.3% at α_2 , 9.4% at α_3 , and 32.8% at α_4 . The underlying analysis attributes the overall pattern to slower, more sluggish mean-reversion dynamics in the atmospheric boundary layer at 100 m relative to the 10 m surface layer; the growing gap at higher lags is not separately explained in that analysis and is reported here as an empirical feature of the fit rather than an interpreted one.

Table 10: CAR(4) coefficients $\alpha_1, \dots, \alpha_4$, both datasets.

Coefficient	Dataset A	Dataset B
α_1	3.3701	3.4676
α_2	4.2546	4.4789
α_3	2.3063	2.5228
α_4	0.3908	0.5188

4.2 Eigenvalues and Stationarity

Stationarity of the continuous-time system requires every eigenvalue of A to have a strictly negative real part, the continuous-time analogue of the discrete unit-circle condition. Table 11 confirms this for both datasets. In both cases the eigenvalues comprise one complex-conjugate pair and two real values; the complex pair generates an oscillatory component in the impulse-response kernel (Section 4.6) with periods that differ substantially across the two datasets: approximately 13.5 days for Dataset A ($2\pi/0.465$) and approximately 24.6 days for Dataset B ($2\pi/0.255$). Both are consistent with mid-latitude synoptic weather cycles, which span a 10–30-day range depending on the atmospheric-circulation pattern. The dominant (least negative) eigenvalue implies a mean-reversion half-life of 2.33 days for Dataset A, close to the 2.0-day half-life already found directly from the discrete AR(4) roots in Section 2.2. For Dataset B the dominant eigenvalue implies a half-life of 1.24 days, markedly faster than the roughly 40-day half-life of the GARCH volatility *regime* found in Section 3.1. This distinction is worth stating explicitly, since the two numbers answer different questions: the CAR(4) eigenvalue governs how quickly the *mean level* of a wind-speed anomaly reverts, while the GARCH persistence governs how quickly the *variance regime* it occurred in reverts, and there is no reason to expect the two rates to coincide.

Table 11: Eigenvalues of the CAR(4) companion matrix, both datasets.

	Dataset A	Dataset B
Complex pair $\lambda_c, \bar{\lambda}_c$	$-0.9358 \pm 0.4653i$	$-1.0730 \pm 0.2550i$
Oscillatory period (days)	13.5	24.6
Real eigenvalue	-1.2004	-0.7614
Dominant real $\lambda_{\min}$	-0.2980	-0.5602
All $\text{Re}(\lambda) < 0$?	Yes	Yes
Half-life, dominant mode (days)	2.33	1.24

4.3 Conditional Forecasts

Under the physical measure, the conditional mean and variance of the state vector at horizon τ are $\mathbb{E}[X(t_0 + \tau) \mid \mathcal{F}_{t_0}] = e_1' \exp(A\tau)X(t_0)$ and $\text{Var}[X(t_0 + \tau) \mid \mathcal{F}_{t_0}] = \sigma_{\text{total}}^2(t_0) \int_0^\tau g(s)^2 ds$, where $g(s) = e_1' \exp(As)e_4$ is the impulse-response kernel and $\sigma_{\text{total}}^2(t_0)$ is the two-component variance of Section 3.2. Both datasets show a non-monotone conditional-mean path, a direct consequence of the oscillatory complex eigenpair: the forecast mean peaks a few days ahead of t_0 before decaying toward zero (Dataset A: peak at a 3-day horizon; Dataset B: peak at a 2-day horizon), rather than declining monotonically from the current state as a simple AR(1)/OU model would predict. Conditional variance saturates at its unconditional level within three to four weeks in both datasets (Dataset A: by a 21–30-day horizon; Dataset B: by a 14–21-day horizon), consistent with the half-lives of Table 11. Table 12 reports the full term structure for both datasets, computed from the state vectors $X(t_0)_A = (0.1299, 0.5159, 0.4672, -0.3422)'$ and $X(t_0)_B = (1.3785, 1.4862, 0.4380, -1.3453)'$, the last four standardised observations of each sample.

Table 12: CAR(4) conditional-mean forecast term structure, both datasets.

Horizon (days)	Dataset A		Dataset B	
	Mean	Std. dev.	Mean	Std. dev.
1	0.779	0.009	2.845	0.031
2	1.303	0.052	3.610	0.172
3	1.460	0.110	3.497	0.358
5	1.113	0.197	2.146	0.614
7	0.658	0.232	0.983	0.699
10	0.270	0.245	0.237	0.719
14	0.082	0.247	0.029	0.720
21	0.010	0.247	0.001	0.720
30	0.001	0.247	≈ 0	0.720

4.4 Normal-Inverse-Gaussian Tails

Section 3.3 found that GARCH-standardised residuals depart from normality in both datasets, in opposite directions: leptokurtic for Dataset A, platykurtic for Dataset B. This motivates fitting a Normal-Inverse-Gaussian (NIG) distribution to the standardised residuals as a heavy-tailed alternative to the Gaussian innovation assumed elsewhere in this chapter, following Barndorff-Nielsen and Shephard [2001]. Table 13 reports the fit. For Dataset A the NIG tail-heaviness parameter $\hat{\alpha} = 4.374$ is small (heavy tails) and the fit is a modest but genuine improvement over the Gaussian (Kolmogorov–Smirnov $p = 0.120$ versus 0.079), though it still falls short of the full empirical kurtosis, an outstanding gap attributed largely to the single 28 June 2017 extreme observation identified in Section 3.3. For Dataset B the fitted $\hat{\alpha} \approx 2,650$ is enormous, reflecting a body so close to Gaussian that the NIG fit offers no material improvement over the normal (KS $p = 0.0275$ versus 0.0269 for the Gaussian). Both fits give a skewness parameter $\hat{\beta}$ close to zero (Dataset A: 0.0067, freely estimated rather than constrained; Dataset

B: -0.0021), consistent with approximately symmetric innovations in both datasets. Given this asymmetry in how much the NIG assumption actually buys, the Gaussian CAR(4) specification is retained as the primary pricing track for Chapter 3 in both datasets, consistent with Benth and Šaltytė Benth [2009]; the NIG fit is carried forward as a documented robustness extension, using the closed-form NIG moment-generating function of Brockwell and Marquardt [2005], rather than as the operative model.

Table 13: Normal-Inverse-Gaussian fit to standardised GARCH residuals, both datasets.

Parameter	Dataset A	Dataset B
$\hat{\alpha}$ (tail heaviness)	4.374	2,649.57
$\hat{\beta}$ (asymmetry)	0.0067	-0.0021
$\hat{\delta}$ (scale)	2.087	51.516
$\hat{\mu}$ (location)	0.000	0.000
Empirical excess kurtosis of $z(t)$	+1.116	-0.278
KS p (NIG / Gaussian)	0.120 / 0.079	0.0275 / 0.0269




Figure 5: CAR(4) structure and conditional forecasts, both datasets. Each panel is arranged in three rows. **Row 1:** three panels – the eigenvalues of the companion matrix A plotted in the complex plane (all in the open left half-plane, confirming stationarity); the impulse-response kernel $g(u) = e_1' \exp(Au)e_4$ against the monotone CAR(1)/Ornstein–Uhlenbeck baseline, showing the local maximum near a 3.6-day horizon that drives the Samuelson-effect violation of Section 4.6; and the squared kernel $g(u)^2$. **Row 2:** two panels – the $[\exp(Au)]_{11}$ mean-propagation curve, and the conditional-forecast fan (mean, with 68% and 95% bands) built forward from the current state $X(t_0)$. **Row 3:** three panels – the conditional standard deviation against horizon, saturating at its unconditional level; a comparison of the CAR(4) and CAR(1) variance term structures; and the full first-row impulse-response matrix $[\exp(Au)]_{1j}$ across all four state components.

4.5 Seasonal Variance Carried Forward

The Fourier seasonal-variance constant C_0 used throughout this chapter and in Chapter 3 is a Phase 1 result (Section 1.3), not re-estimated at this stage: Dataset B’s Phase 4 output carries C_0, C_1, C_2 forward explicitly alongside the CAR(4) parameters, reproducing $C_0 = 1.163$ to the precision of the frozen anchor, while Dataset A’s corresponding output records only the point value $\sigma_{\text{seasonal}}^2(t_0) = 0.2239$ and does not separately re-list C_0 . Both are consistent with $C_0 = 0.13138$; the value is simply carried through from Section 1.3 rather than recomputed, and this thesis does not treat its absence from Dataset A’s Phase 4 output as a gap, since CAR(4) has no reason to refit a quantity that Phase 1 has already fixed.

4.6 The Samuelson Effect

A single-factor Ornstein–Uhlenbeck model ($p = 1$, the Schwartz [1997] special case) has an impulse-response kernel $g(u) = e^{-\alpha_1 u}$ that decays monotonically, and futures volatility built on that kernel therefore rises monotonically as maturity approaches: the classical Samuelson effect. The complex eigenvalue pair found in both datasets’ CAR(4) systems (Table 11) breaks this monotonicity: the kernel $g(u)$ attains a local maximum at a horizon of roughly 3.6 days in both datasets before declining, so that shocks occurring three to four days before contract settlement have *more*, not less, impact on terminal volatility than shocks occurring closer to settlement. Benth and Šaltytė Benth [2009] derive this violation analytically for Nordix wind futures; the same qualitative pattern is confirmed numerically here for both Dataset A and Dataset B, despite their different fitted coefficients, which suggests the effect is a structural property of the CAR($p \geq 2$) kernel rather than an artefact of either specific site.

5 Stochastic-Volatility Framework

The GARCH(1,1) filter of Section 3 is a valid discrete-time description of conditional heteroskedasticity, but it cannot serve directly as the volatility driver of a continuous-time pricing model. This section embeds it into a Heston-type [Heston, 1993] stochastic-volatility system, following the same construction for both datasets; the treatment is theoretical throughout, for reasons explained in Section 5.7, and no parameters are estimated beyond the GARCH proxy mapping.

5.1 Motivation

Two obstacles prevent the GARCH recursion $h(t) = \omega + \alpha\hat{\eta}^2(t-1) + \beta h(t-1)$ from being used directly in a continuous-time pricing framework. First, the Feynman–Kac theorem links a derivative’s price to an expectation under a diffusion’s dynamics precisely because that diffusion possesses a well-defined infinitesimal generator, the operator that encodes its instantaneous drift and diffusion; a GARCH path, updated once per day by a deterministic rule with no diffusion term of its own, has zero quadratic variation

between updates and is therefore not a semimartingale, so no such generator exists and the Feynman–Kac correspondence between a pricing PDE and an expectation simply has nothing to act on. Second, a formal change of measure from the physical measure P to a pricing measure Q , in the sense of Girsanov’s theorem, works by shifting the drift of a Brownian motion that drives the process being repriced; the GARCH filter has no such driving Brownian motion to shift; it is a deterministic function of past innovations, so there is no drift for a change of measure to act on. The remedy, following Heston [1993] and, for the specific application to wind speed, the CAR(p) framework of Benth and Šaltytė Benth [2009] and the Lévy-driven continuous-time machinery of Brockwell and Marquardt [2005], is to embed the GARCH filter $h(t)$ into a Cox–Ingersoll–Ross (CIR) diffusion for a continuous-time variance process $V(t)$, which by construction is a semimartingale driven by its own Brownian motion, with the GARCH parameters entering as moment-matching proxies for the CIR parameters.

5.2 Two-SDE System Under P

For analytical tractability, the wind-speed dynamics are first written in a simplified CAR(1) form, with the full CAR(4) state-vector generalisation given afterward:

$$d\tilde{W}(t) = -\alpha_1 \tilde{W}(t) dt + \sqrt{V(t)} dB_1(t), \quad (5)$$

$$dV(t) = \kappa_V (\bar{v} - V(t)) dt + \xi \sqrt{V(t)} dB_2(t), \quad \langle dB_1, dB_2 \rangle_t = \rho dt, \quad (6)$$

where α_1 is the dominant CAR(4) mean-reversion parameter from Section 4 ($\alpha_1 = 3.3701$ for Dataset A, $\alpha_1 = 3.4676$ for Dataset B), κ_V is the mean-reversion speed of volatility, $\bar{v}$ the long-run variance, and ξ the volatility-of-volatility coefficient. The full state-vector generalisation replaces the scalar drift with $dX(t) = A X(t) dt + e_4 \sqrt{V(t)} dB_1(t)$, where A is the 4×4 companion matrix of Section 4 and $e_4 = (0, 0, 0, 1)'$ injects the diffusion term through the newest lag, consistent with the state-space structure of the discretised CAR(4) model.

5.3 Feller Condition

The CIR process for $V(t)$ remains strictly positive almost surely if and only if the Feller condition $2\kappa_V \bar{v} > \xi^2$ holds; if it is violated, the diffusion term can dominate the mean-reverting drift near the origin and $V(t)$ can reach and remain at zero. The GARCH parameters of Section 3.1 are mapped to CIR proxies by moment matching, $\kappa_V \approx 1 - (\alpha + \beta)$, $\bar{v} \approx \omega / (1 - \alpha - \beta)$, and $\xi^2 \approx \text{Var}[h(t)]$, exactly as specified for the numerical Feller check permitted within the scope of this theoretical chapter. For Dataset B, this gives $\kappa_V = 0.0179$, $\bar{v} = 0.7522$, $\xi^2 = 0.000959$, so that $2\kappa_V \bar{v} = 0.0270$ comfortably exceeds ξ^2 : the Feller condition holds with a margin of $28.1\times$.

For Dataset A, the same moment-matching procedure gives $\kappa_V = 0.0165$ and $\bar{v} = 0.608$, in agreement with Dataset B’s computation to four decimal places, but $\xi^2 = \text{Var}[h(t)] = 0.0253$, the sample variance of the full Dataset A $h(t)$ series computed by the same method used for Dataset B. Substituting into the condition gives $2\kappa_V \bar{v} = 0.0201$,

which is *smaller* than ξ^2 : the Feller condition is violated, with a margin of $0.8\times$. Unlike Dataset B, the theoretical CIR embedding of this chapter is therefore not well posed for the proxy dataset on this specific numerical check.

This violation is not a modelling failure so much as a direct symptom of the proxy dataset's site characteristics. Dataset A (10 m) is heavily influenced by boundary-layer friction and local diurnal anomalies, producing the leptokurtic residual distribution and the single extreme shock already identified in Section 3.3 (the 5.79σ event of 28 June 2017). In a GARCH filter with the comparatively high shock sensitivity estimated for Dataset A ($\alpha = 0.0311$, against 0.0059 for Dataset B), an outlier of that size produces a sharp, transient spike in $h(t)$, exactly the kind of feature that inflates a sample-variance-based ξ^2 proxy computed directly from the $h(t)$ path. Dataset B, measured at 100 m in the free atmosphere and free of any comparable shock, produces a markedly smoother $h(t)$ path and comfortably clears the condition instead. The asymmetry is informative in its own right, distinct from and additional to the pricing-error evidence of Chapter 3: the theoretical stochastic-volatility embedding developed in this chapter is well posed for the reference dataset and is not well posed, on this specific check, for the geographic proxy.

5.4 Remark on the Change of Measure

A formal change of measure from P to a pricing measure Q would introduce market-price-of-risk parameters θ_1, θ_2 shifting the drifts of B_1 and B_2 , with $\kappa_V^Q = \kappa_V + \theta_2\xi$ under Q . This thesis adopts $\theta_1 = \theta_2 = 0$ throughout as a benchmark assumption, in the absence of EEX option-implied data, for two reasons that hold identically for both datasets. First, SMARD daily wind generation is a measured physical quantity, not a traded financial asset; the replication argument that pins down a unique risk-neutral measure for a traded derivative does not apply here, since no self-financing hedging portfolio in the underlying exists. Second, the only actively traded wind derivatives, the Nordix futures, reference Norwegian rather than German wind, and the geographic mismatch (Bologna or Nordfriesland versus Norway) prevents any clean calibration of θ from that market even in principle. Under $\theta = 0$, $P = Q$, and Track A (empirical realised variance) and Track B (model-implied variance) are defined on a common measure; the variance-swap fair strike of Chapter 3 reduces to the closed-form Tol [1997] expression $K_{\text{var}}^B = h(t_0) \times C_0$, and no Girsanov computation is required anywhere in this thesis. This benchmark is adopted as a tractable simplification given the absence of option-implied data for this specific underlying, not as a claim that the market price of risk for wind-linked instruments is in fact zero.

5.5 Characteristic Function and Round-Trip Consistency

Under the two-SDE system, the log-characteristic function of $\tilde{W}(T)$ conditional on $\mathcal{F}_t$ takes the Heston [1993] form $\log \phi(u; t, T) = iu \mathbb{E}[\tilde{W}(T) \mid \mathcal{F}_t] + C(\tau; u)\bar{v} + D(\tau; u)V(t)$, $\tau = T - t$, where C and D satisfy a coupled Riccati system, an ordinary differential equation containing a quadratic term in D , arising because the volatility process enters

its own diffusion coefficient through $\sqrt{V(\tau)}$. This is what makes the system harder to solve than the linear CAR(4) dynamics of Section 4, whose companion-matrix ODE is solved exactly by the matrix exponential $\exp(A\tau)$: a Riccati equation has no such generic closed-form solution and instead requires the specific change-of-variables technique that Heston [1993] derives for this model, exploiting its particular algebraic structure rather than a general linear-systems result. Heston’s closed-form solution defines C and D via complex square roots and logarithms, a substantially more delicate object to evaluate numerically than the CAR(4) matrix exponential used elsewhere in this chapter, since a careless implementation can pick the wrong branch of the complex logarithm and return a discontinuous, incorrect path as τ varies.

Under the $\theta = 0$ benchmark, the first moment collapses to $\mathbb{E}[\tilde{W}(T) \mid \mathcal{F}_t] = e_1' \exp(A\tau)X(t)$, exactly the CAR(4) conditional-mean forecast of Section 4. It is worth being precise about what this identity does and does not test: because $\theta = 0$ removes C and D from the first moment entirely, leaving only the linear companion-matrix propagator, the round-trip check that follows verifies the drift structure common to both representations, not the numerical handling of C and D themselves, which enter only the variance and higher moments of the characteristic function and are not separately checked here. This identity is verified numerically for both datasets by comparing the characteristic function’s first moment against the Phase 4 conditional-forecast table at eleven horizons from 1 to 90 days: the maximum absolute discrepancy is 4.44×10^{-16} in both datasets, machine precision, confirming that the stochastic-volatility embedding’s drift term is numerically consistent with the discrete CAR(4) forecasts it is built to extend, independently of the Feller-condition question raised above.

5.6 Variance Term Structure

The CAR(4) kernel $g(s) = e_1' \exp(As)e_4$ satisfies $g(0) = 0$: new information enters only through the most recent of the four lagged state components, and must propagate forward before it reaches the first. As a result, the piecewise-constant- h variance term structure of Tol [1997], $\text{Var}^Q[\tilde{W}(T) \mid \mathcal{F}_t] = h(t_0) \sigma_{\text{seasonal}}^2(t_0) \int_0^\tau g(s)^2 ds$, starts at zero and ramps up slowly. A simplified Heston/CIR term structure built on the CAR(1) reduction of Section 5.2 instead uses a kernel $g(s) = e^{-\alpha_1 s}$ that starts at one, producing an immediate jump in short-horizon variance. For Dataset B, the two term structures differ by a factor of roughly four at a one-day horizon (Tol/CAR(4): 0.00247; Heston/CAR(1): 0.00985) and converge only beyond a roughly two-week horizon, after which the ratio settles in the 0.57–0.93 range through 90 days. The same qualitative pattern holds for Dataset A. This divergence is not a modelling inconsistency between the two approximations; it reflects a genuine structural feature of the CAR(4) state-space representation, and it is a first-order consideration for any option or short-dated instrument priced off the CAR(4) kernel rather than the simplified CAR(1) proxy.




Figure 6: Variance term structure, both datasets: the piecewise-constant- h CAR(4) approximation of Tol [1997] against the simplified Heston/CIR CAR(1) term structure of Section 5.6. Each panel contains two plots, side by side. **Left:** the full term structure over a 0–90-day horizon, showing the two approximations converging at long horizons. **Right:** a zoomed view over 0–14 days, isolating the short-horizon divergence driven by the CAR(4) kernel’s $g(0) = 0$ property.

5.7 Why Maximum-Likelihood Estimation Is Deferred

The CIR parameters $(\kappa_V, \bar{v}, \xi, \rho)$ above are moment-matching proxies from the GARCH fit, not maximum-likelihood estimates of the continuous-time system itself, and this thesis does not attempt that estimation at daily frequency. The reason is an identification problem inherent to slow-mean-reverting variance processes sampled coarsely relative to their own half-life. With $\kappa_V \approx 0.018$ in both datasets, the implied mean-reversion half-life is on the order of 40 days, so that a daily observation interval provides roughly 40–60 observations per e-folding time; separating the mean-reversion speed κ_V from the diffusion coefficient ξ requires enough independent decay cycles to distinguish “reverts slowly with small shocks” from “reverts moderately with occasional large shocks,” and at daily frequency the two are close to observationally equivalent. Dataset A compounds this with its short four-year window (1,462 daily observations, of which only a few dozen independent decay cycles are available) and its 10 m measurement height, which would require a further profile correction before any estimated variance dynamics could be applied to hub-height pricing instruments. The designated upgrade path for both datasets is hourly ECMWF ERA5 reanalysis data at the same locations, which would increase the observation count by roughly two orders of magnitude and make the ξ/κ_V ratio separately identifiable; this is left to future work; for the daily-frequency scope of this thesis, the GARCH-derived proxies under the $\theta = 0$, piecewise-constant- h approximation are the operative tool, and are what Chapter 3 uses throughout.

References

- Alexandridis, A., & Zapranis, A. (2013). Wind derivatives: modeling and pricing. *Computational Economics*, 41(3), 299–326.
- Barndorff-Nielsen, O. E., & Shephard, N. (2001). Non-Gaussian Ornstein–Uhlenbeck-based models and some of their uses in financial economics. *Journal of the Royal Statistical Society: Series B*, 63(2), 167–241.
- Benth, F. E., & Šaltytė Benth, J. (2009). Dynamic pricing of wind futures. *Energy Economics*, 31(1), 16–24.
- Bollerslev, T. (1986). Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3), 307–327.
- Brockwell, P. J., & Marquardt, T. (2005). Lévy-driven and fractionally integrated ARMA processes with continuous time parameter. *Statistica Sinica*, 15(2), 477–494.
- Diebold, F. X., & Mariano, R. S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263.
- Engle, R. F. (1982). Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4), 987–1007.
- Garcia, A., Torres, J. L., Prieto, E., & de Francisco, A. (1998). Fitting wind speed distributions: a case study. *Solar Energy*, 62(2), 139–144.
- Haslett, J., & Raftery, A. E. (1989). Space-time modelling with long-memory dependence: assessing Ireland’s wind power resource. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 38(1), 1–50.
- Heston, S. L. (1993). A closed-form solution for options with stochastic volatility with applications to bond and currency options. *Review of Financial Studies*, 6(2), 327–343.
- Kavasseri, R. G., & Seetharaman, K. (2009). Day-ahead wind speed forecasting using f-ARIMA models. *Renewable Energy*, 34(5), 1388–1393.
- Schwartz, E. S. (1997). The stochastic behavior of commodity prices: implications for valuation and hedging. *Journal of Finance*, 52(3), 923–973.
- Tol, R. S. J. (1997). Autoregressive conditional heteroscedasticity in daily wind speed measurements. *Theoretical and Applied Climatology*, 56(1–2), 113–122.
- Tuller, S. E., & Brett, A. C. (1984). The characteristics of wind velocity that favor the Weibull distribution in wind speed analysis. *Journal of Climate and Applied Meteorology*, 23(1), 124–134.